

Rebekah Warner

Fairfax, VA | rebekahwarner@protonmail.com

SKILLS

Security Clearance: TS/SCI w/ FS Poly

Languages: Python, R, SQL (SQLite, MySQL, PostgreSQL), HTML, CSS

Machine Learning: Regression (linear/logistic), random forests, SVM, k-means, KNN, Naive Bayes; model evaluation; scikit-learn

NLP: Text preprocessing, TF-IDF, sentiment/text classification, topic modeling (LDA), NLTK

Generative AI: LLM APIs (OpenAI/Anthropic), prompt engineering

Python Libraries: pandas, NumPy, SciPy, scikit-learn, NLTK, Matplotlib, Seaborn, Plotly, WordCloud, BeautifulSoup, Requests, SQLAlchemy, psycpg2, joblib, openpyxl, Streamlit

Data Engineering & Cloud: Database design, ETL, data cleaning (pandas, regex), AWS, Docker, Git/GitHub, Jupyter

Visualization & BI: Tableau, Power BI, Matplotlib, Seaborn, Plotly

Familiar with: PyTorch, A/B testing / controlled experiments (coursework)

Foreign Language: Chinese (professional proficiency — speaking, reading, writing)

EDUCATION

University of Maryland — College Park, MD

B.S., Social Data Science: Government and Politics Track

EXPERIENCE

Guidehouse — Data Scientist

Tysons, VA | Aug 2025 – Present

- Designed and optimized database architectures to support advanced analytics and data science workflows.
- Built an automated schema-drift monitor that snapshots all tables, views, and columns of the production database to a JSON baseline, diffs the live schema on each run, and reports changes to Excel — converting a reactive failure mode into proactive detection.
- Piloted migration of data-processing workflows to an AWS environment; automated scripts within a locked-down non-admin setting and advocated a Linux transition to expand team capabilities.
- Reduced per-person environment setup from roughly three days to about 15 minutes by building batch-file automation and a supporting test suite.
- Fulfilled high-volume ad-hoc data requests with custom SQL on tight turnaround; maintained and debugged production dashboards.
- Mentored new team members on coding fundamentals and regularly communicated technical work to non-technical stakeholders.
- Conducted exploratory data-quality analysis across large populations and translated findings into an executive PowerPoint briefing delivered to the client.

Red Alpha — Data Science Intern

Columbia, MD | Jun 2024 – Aug 2024

- Built and deployed a production capacity-planning application (Python, PostgreSQL, Streamlit) with two interns, enabling clients to optimize resource allocation and saving hundreds of hours.
- Developed dynamic PostgreSQL-driven visualizations giving leadership clear insight into the project pipeline.
- Tuned database performance through benchmarking, query optimization, and indexing; cleaned data with regex to ensure quality.
- Identified team inefficiencies and advocated a sprint system that improved project management and productivity.

U.S. Department of Defense — Language Analyst

Sep 2020 – Aug 2023 | Ft. Meade, MD

- Performed regular data analysis and pattern-finding; built Jupyter notebooks that consolidated large datasets for more efficient workflows.
- Completed hundreds of translations informing 100+ intelligence reports reviewed by senior U.S. government officials.
- Collaborated across diverse teams to achieve mission goals.

PROJECTS

NLP Sentiment Text Classifier — Python, scikit-learn, NLTK

- Built an end-to-end binary sentiment classifier on 10,000 labeled text documents using NLTK preprocessing (tokenization, stopword removal; joblib-parallelized), TF-IDF features (unigrams and bigrams), and a linear-kernel SVM; reached 87% validation accuracy on an 80/20 split and generated predictions on a 5,000-document test set.

LinkedIn IT Job-Market Analysis — Python, pandas, seaborn

- Merged and cleaned eight large LinkedIn datasets (60+ MB job-postings file) into a unified table via multiple keyed joins; normalized inconsistent salary data to a common yearly basis, handled missing values, and produced exploratory visualizations on the U.S. IT & Services job market.